

Prakash Sharma

AI/ML & Agentic AI Developer

prakash.sharma.1462003@gmail.com | +91 8349353477 | Pune, India (Open to Remote / Relocate)

[linkedin.com/in/prakash-sharma-ps](https://www.linkedin.com/in/prakash-sharma-ps) | github.com/light14

PROFESSIONAL SUMMARY

AI/ML Developer with production experience building Agentic AI applications, enterprise RAG pipelines, and Python microservices. Hands-on experience with multi-agent workflow orchestration, vector databases (Milvus), LLM APIs, and machine learning models (Scikit-learn, Pandas). Reduced end-to-end RAG response latency from 30+ seconds to under 5 seconds. Owned features from architecture through prompt engineering, testing, and production delivery.

TECHNICAL SKILLS

AI & Agentic AI: Agentic AI Workflows, Multi-Agent Systems, LlamaIndex, LangChain concepts, RAG, Prompt Engineering, Tool Calling, OpenAI API, Semantic Search

Machine Learning & Data: NumPy, Pandas, Scikit-learn, ML Model Training, Feature Engineering, Data Cleaning

Backend & APIs: Python, FastAPI, Pydantic, REST APIs, AsyncIO, Microservices

Storage & Databases: Milvus (Vector DB), MongoDB, GridFS, PostgreSQL, Metadata Isolation

Tools & Automation: Git, GitHub, PyMuPDF4LLM, Azure OCR, Web Scraping (Playwright, Selenium)

PROFESSIONAL EXPERIENCE

AI & Data Engineer | Tradxlink Private Limited | Remote (Noida-based) | Sep 2025 – Present

- Architected a production multi-agent system using agent workflow frameworks for query decomposition, context routing, API tool calling, and response synthesis in Python.
- Built enterprise RAG ingestion pipelines in Milvus using metadata filtering, public/private context isolation, and prompt engineering to reduce LLM hallucinations.
- Cut end-to-end RAG response latency from 30+ seconds to under 5 seconds by replacing multi-step tree summaries with compact generation flows and optimizing API roundtrips.
- Developed Python REST microservices using FastAPI for voyage delay prediction, vessel reliability scoring, and automated document feature extraction.
- Resolved large-document index failures by redesigning query-engine persistence around rebuildable GridFS references instead of in-memory graph objects.
- Built document extraction pipelines combining PyMuPDF4LLM and Azure OCR to transform multi-page trade PDFs into structured JSON schemas.

Process Executive Intern (Data) | Augtech NextWealth | On-site | Jun 2025 – Sep 2025

- Annotated LiDAR data for AV object detection/scene understanding, sustaining 95–97% quality and high precision against QA benchmarks.

Operations Intern | Edxso | Remote (Noida-based) | Apr 2025 – Jun 2025

- Automated web data extraction with Python + Selenium and Google Sheets API, cutting manual collection effort by ~80% and improving reporting speed.

PROJECTS

Mira — AI Receptionist (Beta/Testing) | Python, LLMs, FastAPI, Telegram/WhatsApp | **Demo:** <https://mira.voxlabs.site/>

- Developed an agentic booking receptionist for service businesses featuring conversational scheduling, context management, and customer onboarding via messaging APIs.

Redax — AI Redaction Security Auditor (In Progress) | Python, OCR, LLMs, Document Intelligence

- Built a document auditor that scans redacted PDFs using OCR and LLMs to identify residual PII and secret leaks, generating automated risk reports.

Credit Card Fraud Detection | Python, Scikit-learn, Pandas, Streamlit

- Built a supervised ML classification application using Random Forest models to detect fraudulent transactions from imbalanced financial datasets.

EDUCATION

Bachelor of Computer Applications (BCA) | Rungta College of Science and Technology, Chhattisgarh | CGPA: 8.0 | 2020 – 2023

CERTIFICATIONS

- Python (GUVI, Aug 2023)
- Data Science Internship (Internship Studio, Aug 2024 – Feb 2025)